Clustering, k-means, k-means++ and the advantages of careful seeding

- David Arthur, Sergei Vassilvitskii. *k-means++: The Advantages of Careful Seeding*. In SODA 2007

What is clustering?

- a grouping of data objects such that the objects within a group are similar (or near) to one another and dissimilar (or far) from the objects in other

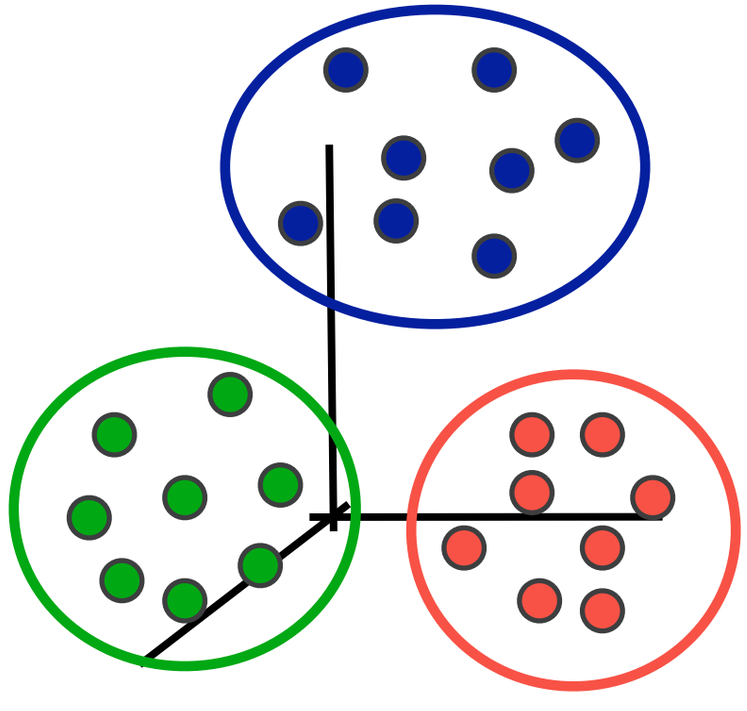

How to capture this objective?

a grouping of data objects such that the objects within a group are similar (or near) to one another and dissimilar (or far) from the objects in other groups

minimize intra-cluster distances

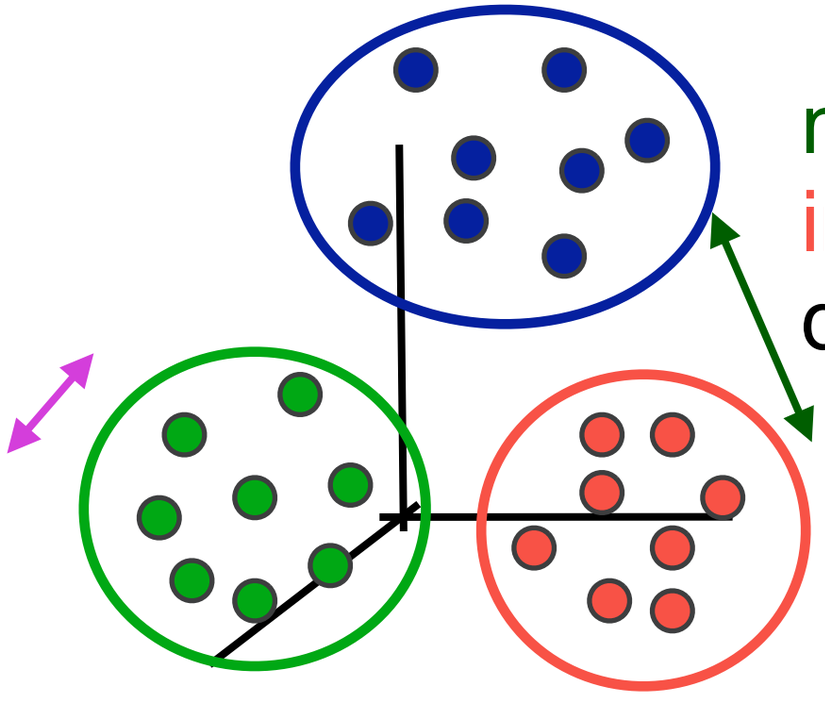

maximize inter-cluster distances

The clustering problem

- Given a collection of data objects
- Find a grouping so that
  - similar objects are in the same cluster
  - dissimilar objects are in different clusters

- Why we care ?

- stand-alone tool to gain insight into the data
  - visualization
- preprocessing step for other algorithms
  - indexing or compression often relies on clustering

Applications of clustering

- image processing
  - cluster images based on their visual content
- web mining
  - cluster groups of users based on their access patterns on webpages
  - cluster webpages based on their content
- bioinformatics
  - cluster similar proteins together (similarity wrt chemical structure and/or functionality etc)
- many more...

The clustering problem

- Given a collection of data objects
- Find a grouping so that
  - similar objects are in the same cluster
  - dissimilar objects are in different clusters

- Basic questions:
  - what does similar mean?
  - what is a good partition of the objects?

    i.e., how is the quality of a solution measured?
  - how to find a good partition?

Notion of a cluster can be ambiguous

How many clusters?

Six Clusters

Two Clusters

Four Clusters

Types of clusterings

- Partitional
  - each object belongs in exactly one cluster
- Hierarchical
  - a set of nested clusters organized in a tree
- Density Based
  - clustering is defined based on the local density of the points

Partitional clustering

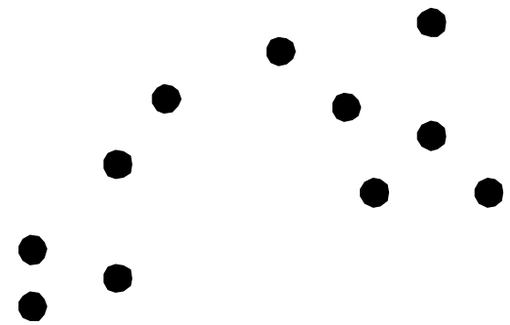

Original Points

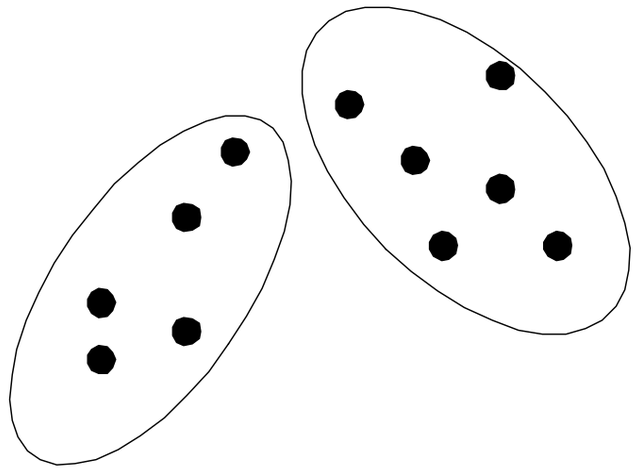

A Partitional Clustering

Partitional algorithms

- partition the n objects into k clusters
  - each object belongs to exactly one cluster
  - the number of clusters k is given in advance

The k-means problem

- consider set $X=\{x_1,...,x_n\}$ of n points in R^d
- assume that the number k is given
- problem:
  - find k points $c_1,...,c_k$ (named centers or means)

so that the cost

$$\sum_{i=1}^{n} \min_j \left\{ L_2^2(x_i, c_j) \right\} = \sum_{i=1}^{n} \min_j ||x_i - c_j||_2^2$$

is minimized

The k-means problem

- consider set $X=\{x_1,...,x_n\}$ of n points in R^d
- assume that the number k is given
- problem:
  - find k points $c_1,...,c_k$ (named centers or means)
  - and partition X into $\{X_1,...,X_k\}$ by assigning each point x_i in X to its nearest cluster center,
  - so that the cost

$$\sum_{i=1}^{n} \min_j ||x_i - c_j||_2^2 = \sum_{j=1}^{k} \sum_{x \in X_j} ||x - c_j||_2^2$$

  is minimized

The k-means problem

- k=1 and k=n are easy special cases (why?)
- an NP-hard problem if the dimension of the data is at least 2 (d≥2)
- in practice, a simple iterative (greedy) algorithm works quite well

The k-means algorithm

- voted among the top-10 algorithms in data mining
- one way of solving the k-means problem

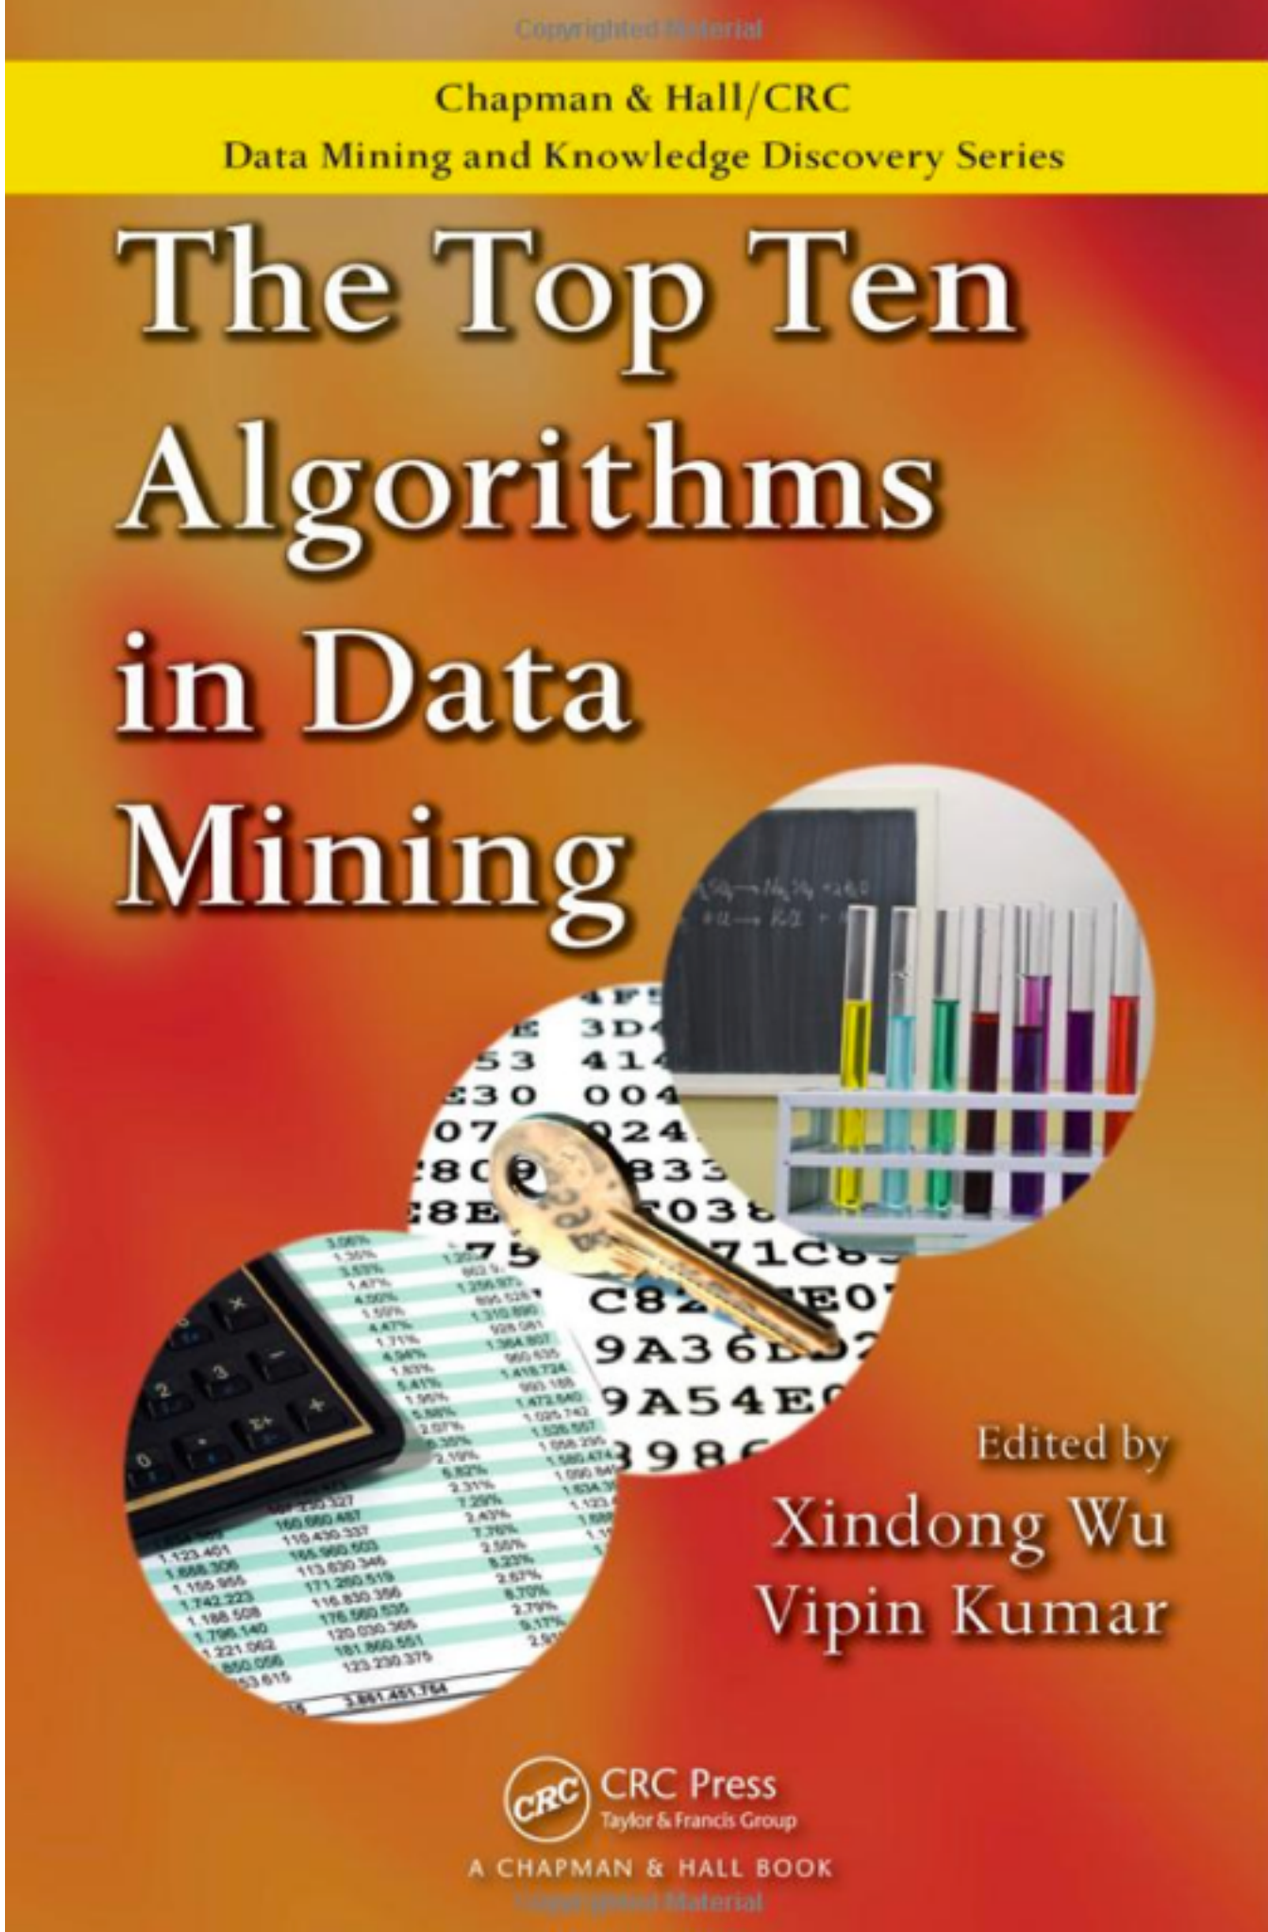

The k-means (Lloyd's) algorithm

1. randomly (or with any other method) pick k cluster centers $\{c_1,\ldots,c_k\}$
2. for each j, set the cluster X_j to be the set of points in X that are the closest to center c_j
3. for each j let c_j be the center of mass of cluster X_j (mean of the vectors in X_j)
4. repeat (go to step 2) until convergence

Sample execution

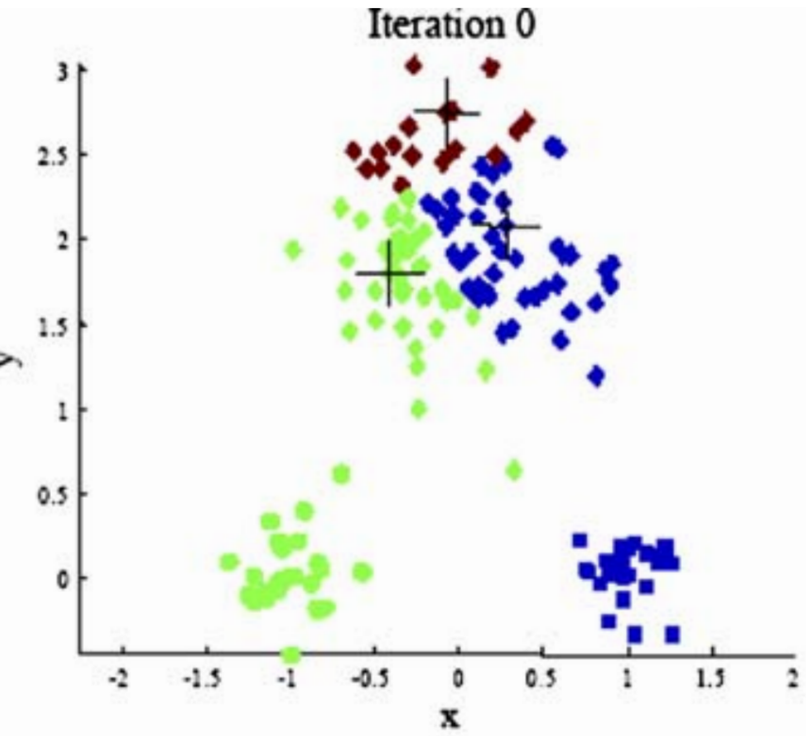

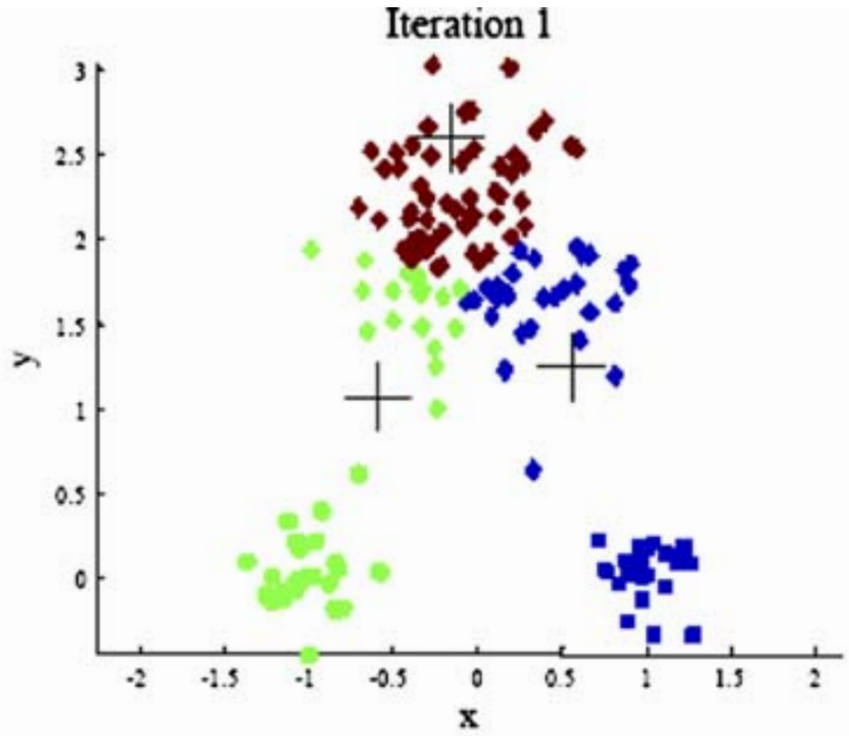

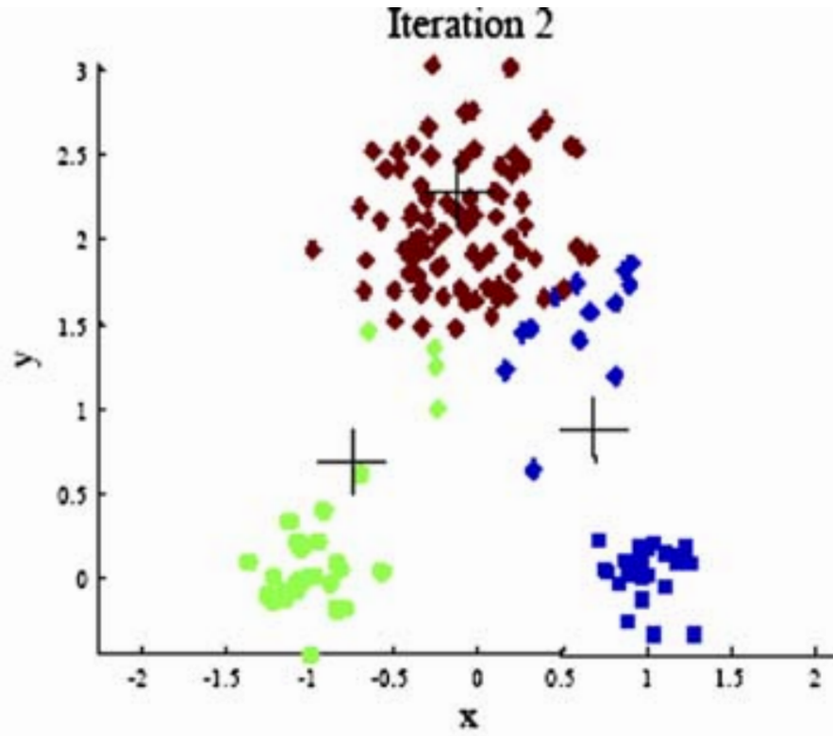

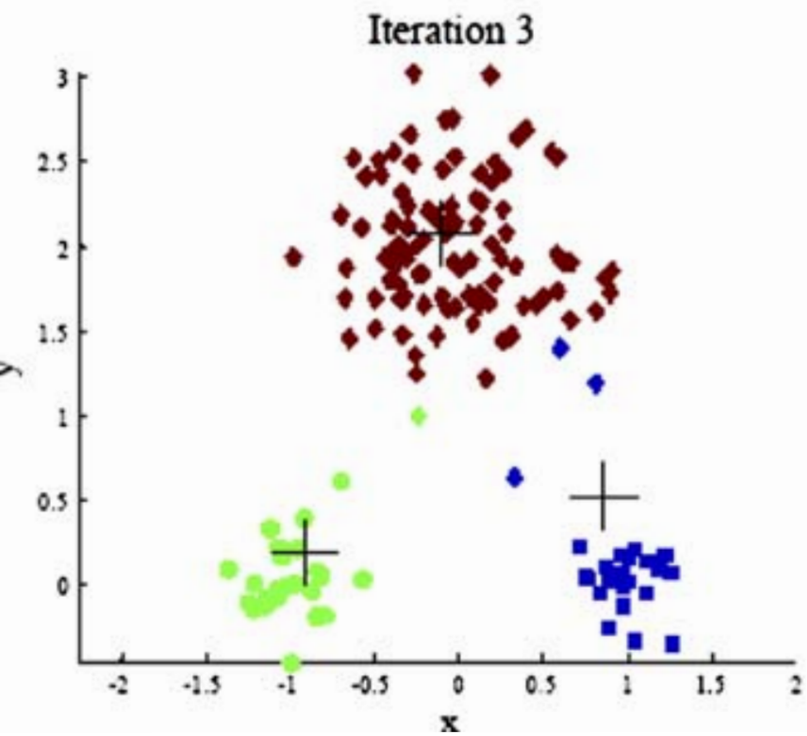

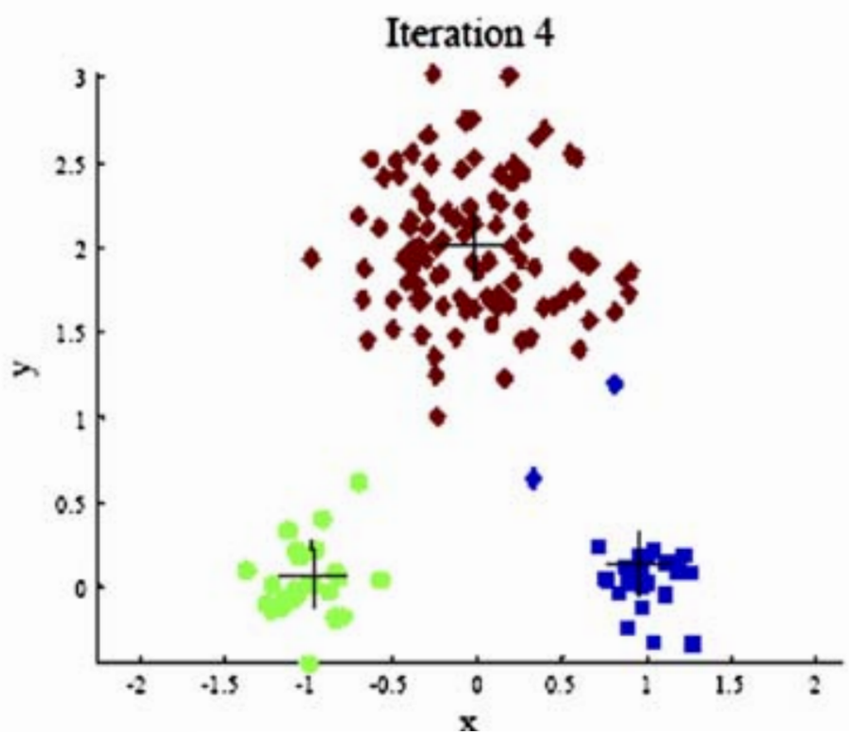

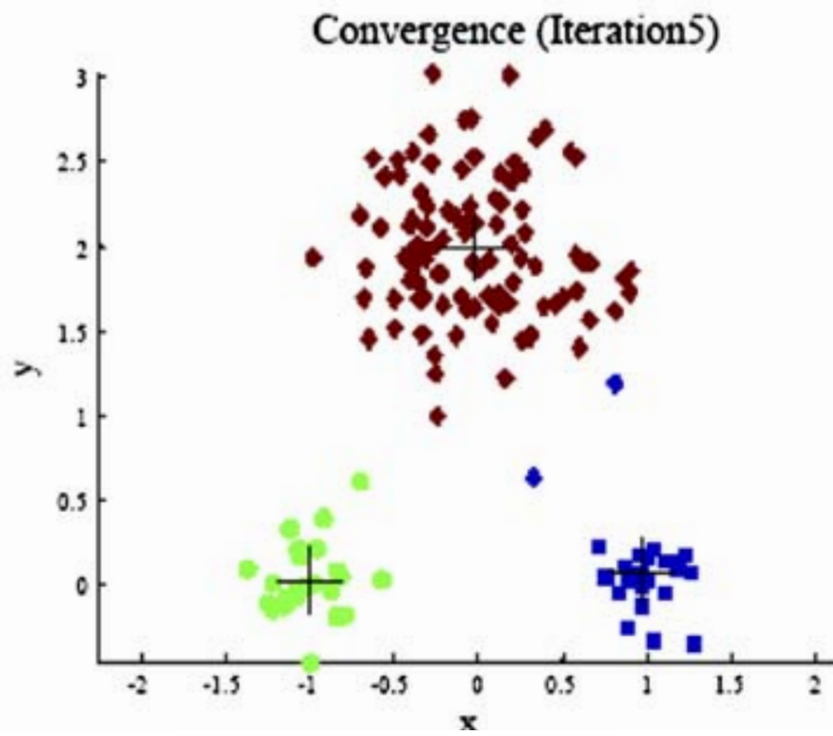

Properties of the k-means algorithm

- finds a local optimum
- often converges quickly
  - but not always
- the choice of initial points can have large influence in the result

Effects of bad initialization

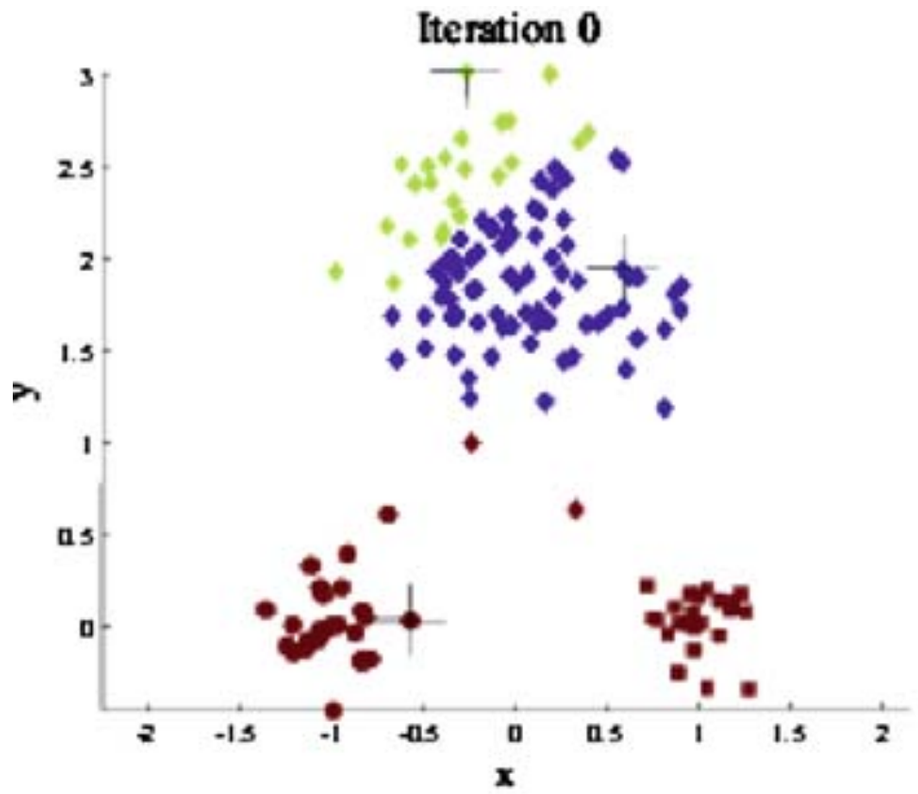

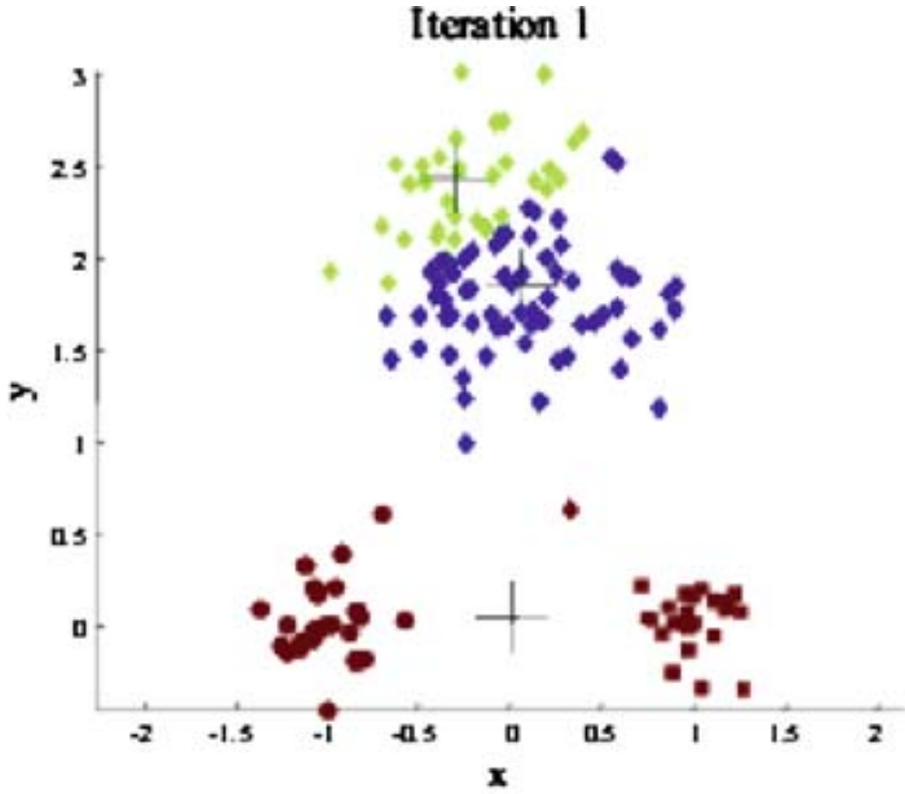

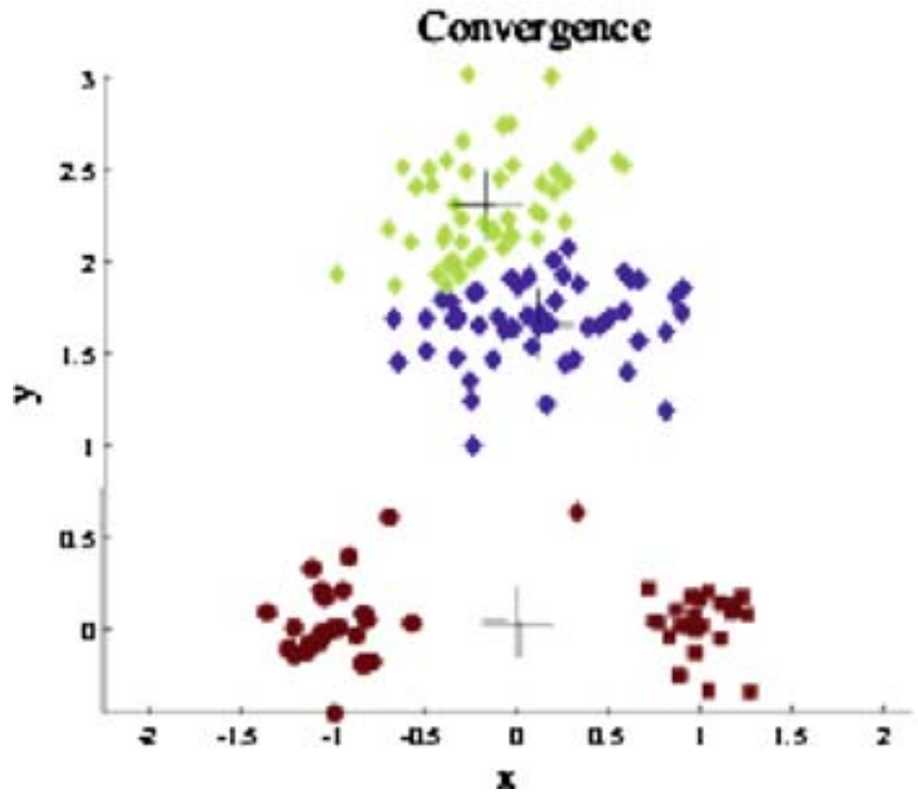

Limitations of k-means: different sizes

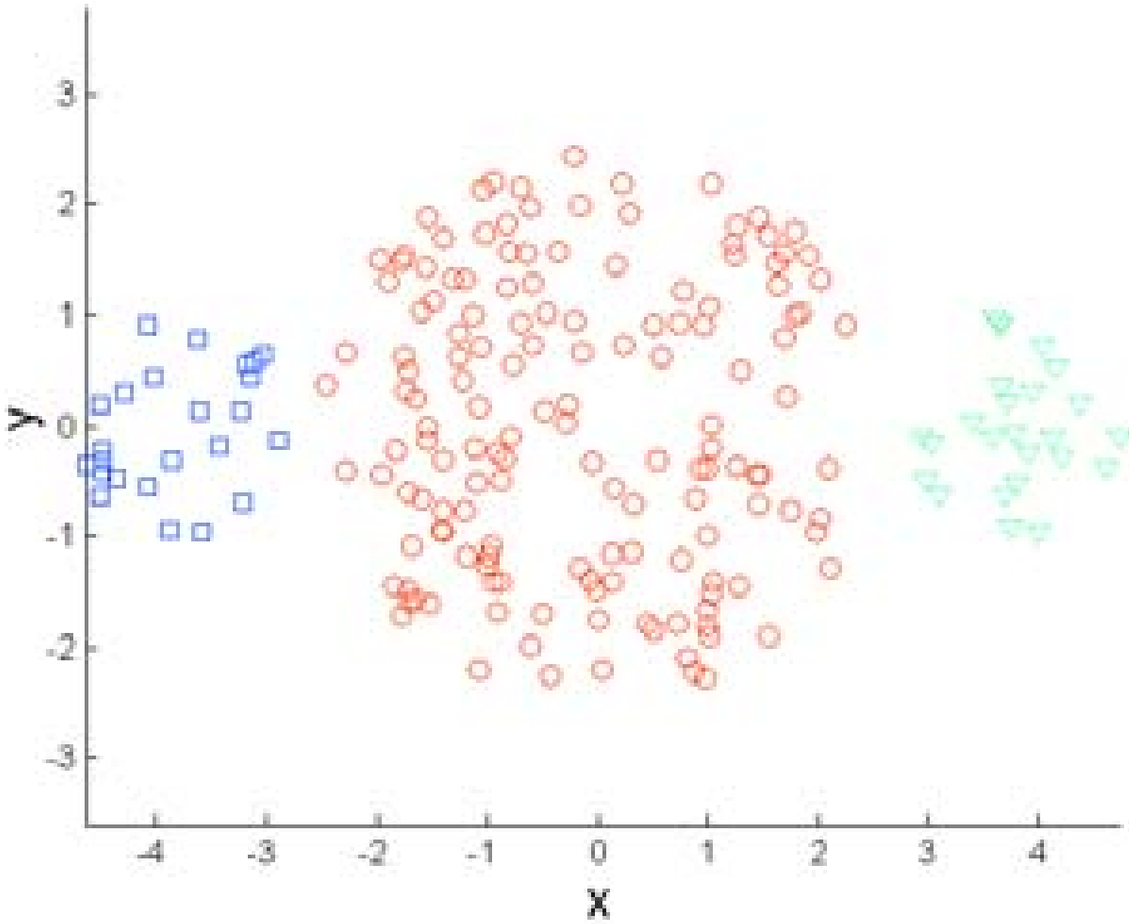

Original Points

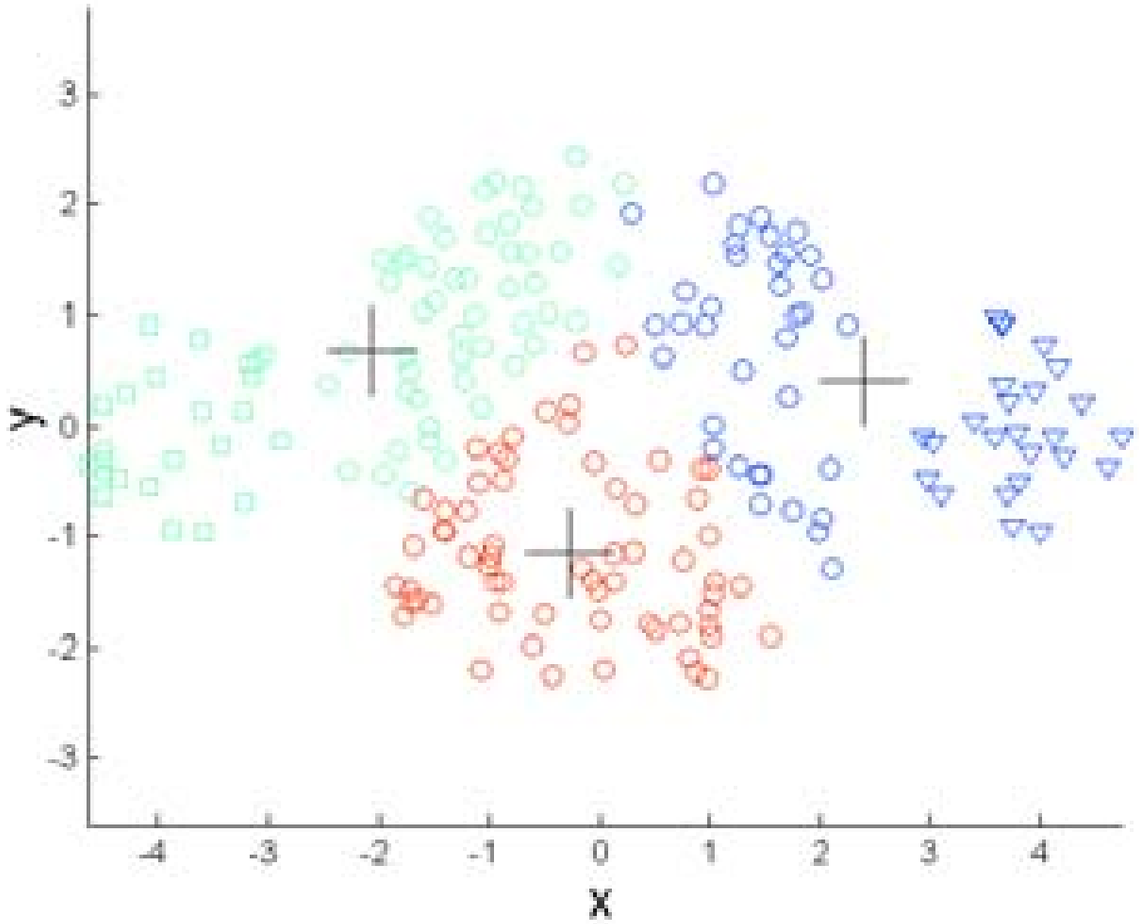

K-means (3 Clusters)

Limitations of k-means: different density

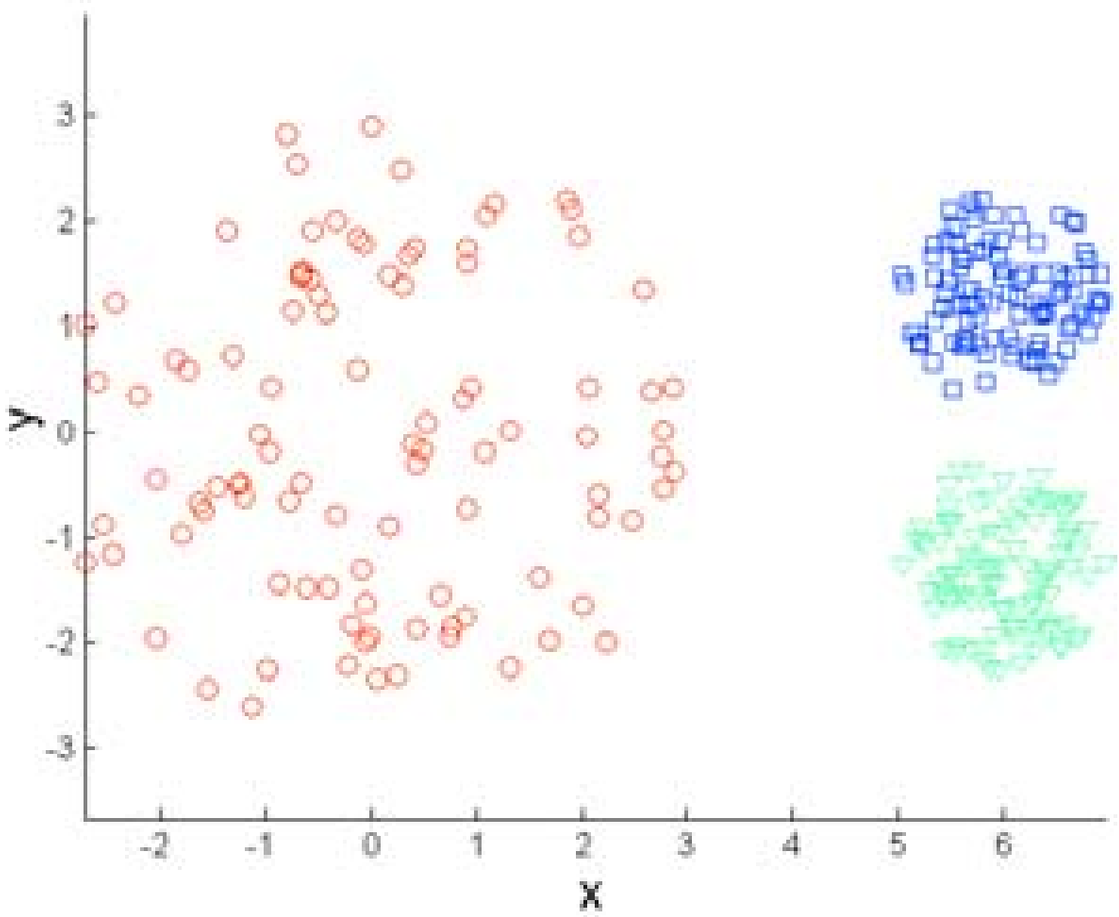

Original Points

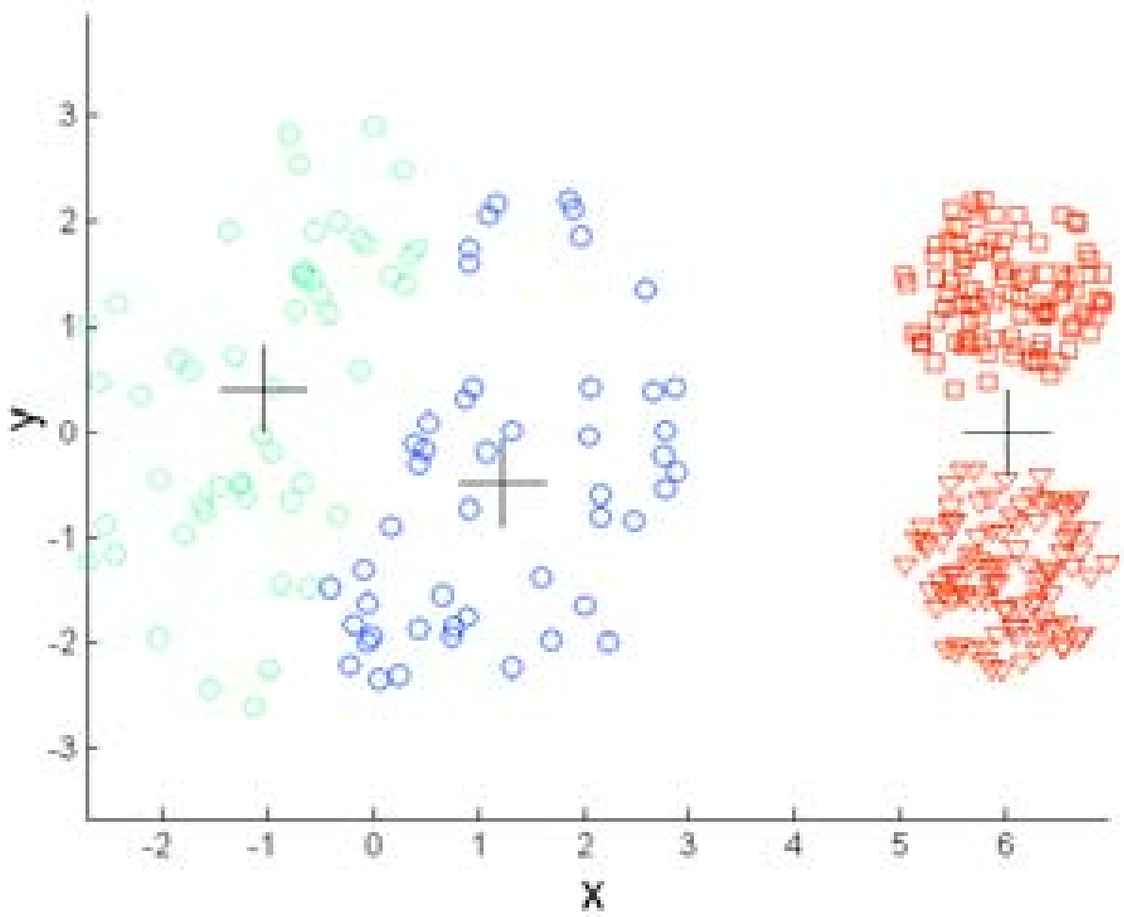

K-means (3 Clusters)

Limitations of k-means: non-spherical shapes

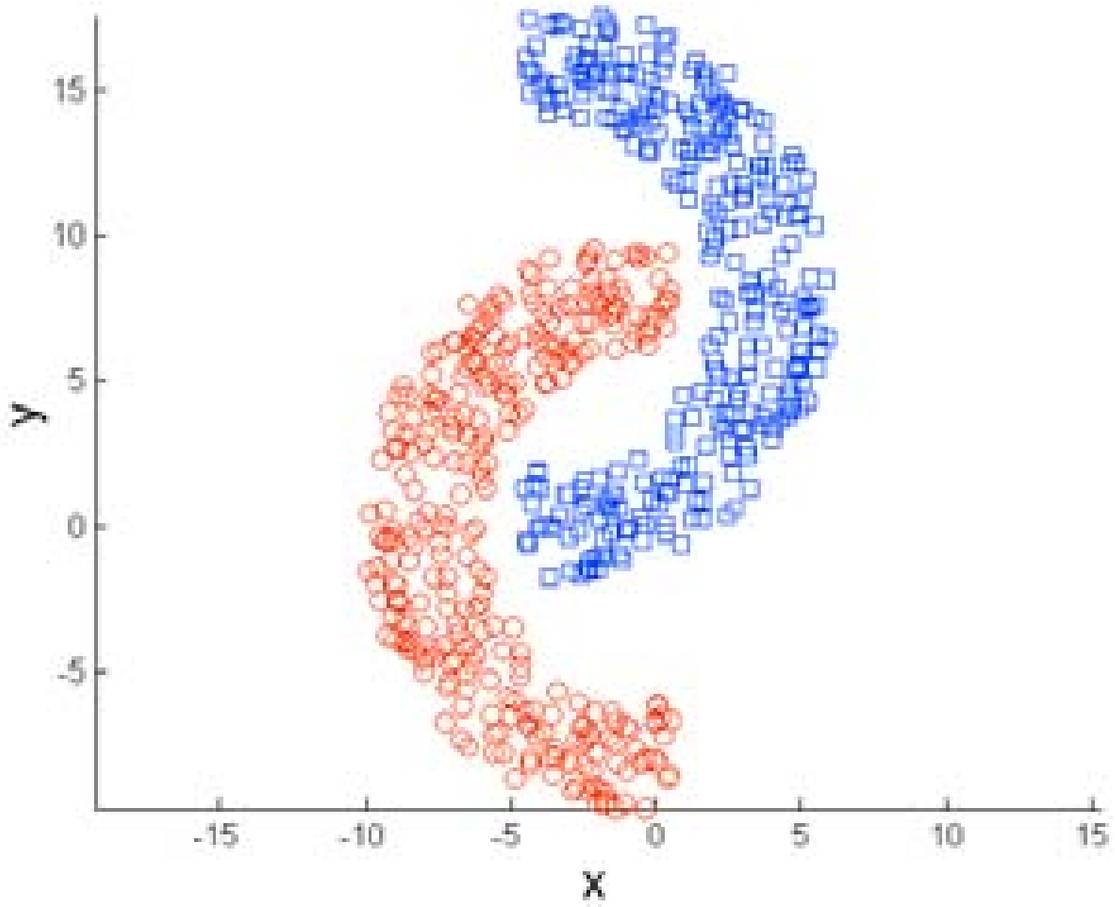

Original Points

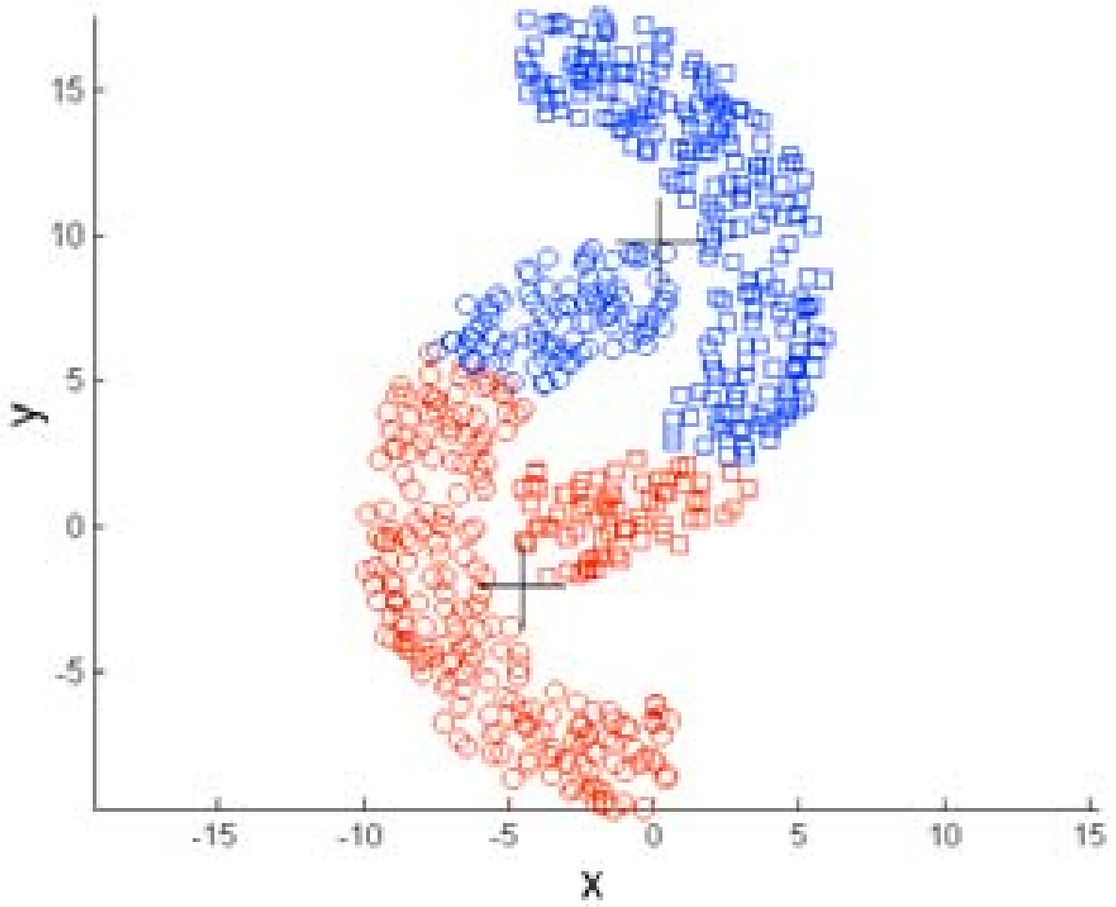

K-means (2 Clusters)

Discussion on the k-means algorithm

- finds a local optimum
- often converges quickly
  but not always
- the choice of initial points can have large influence in the result
- tends to find spherical clusters
- outliers can cause a problem
- different densities may cause a problem

Initialization

- random initialization
- random, but repeat many times and take the best solution
  - helps, but solution can still be bad
- pick points that are distant to each other
  - k-means++
  - provable guarantees

k-means++

David Arthur and Sergei Vassilvitskii

k-means++: The advantages of careful seeding

SODA 2007

k-means algorithm: random initialization

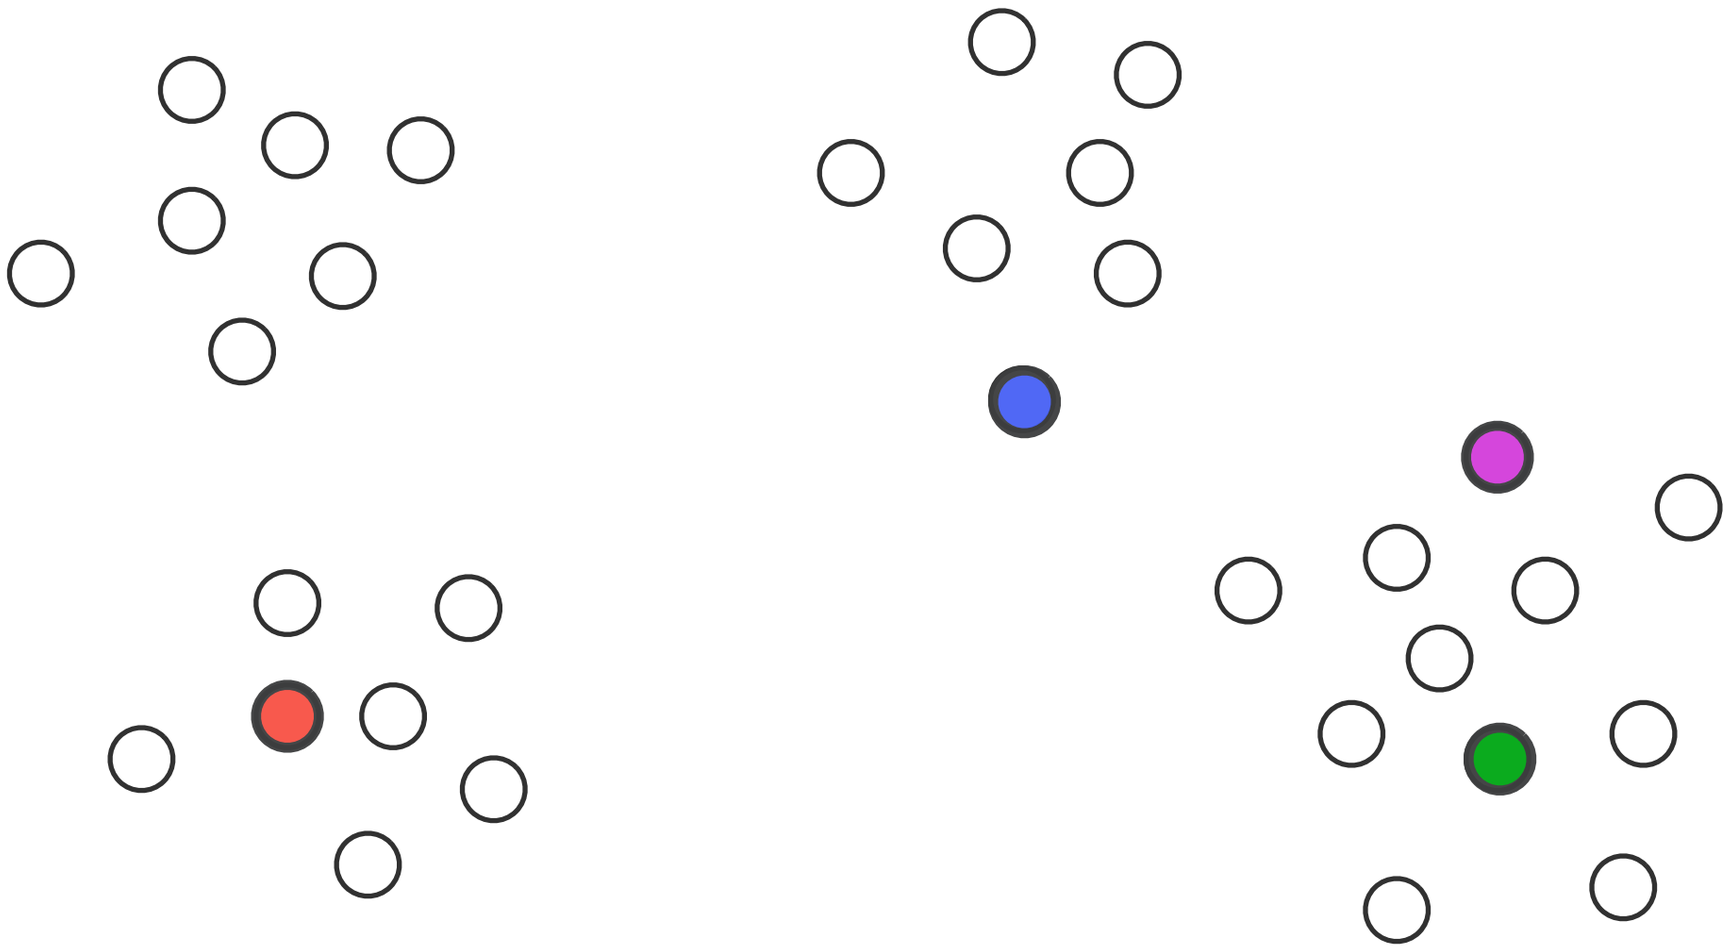

k-means algorithm: random initialization

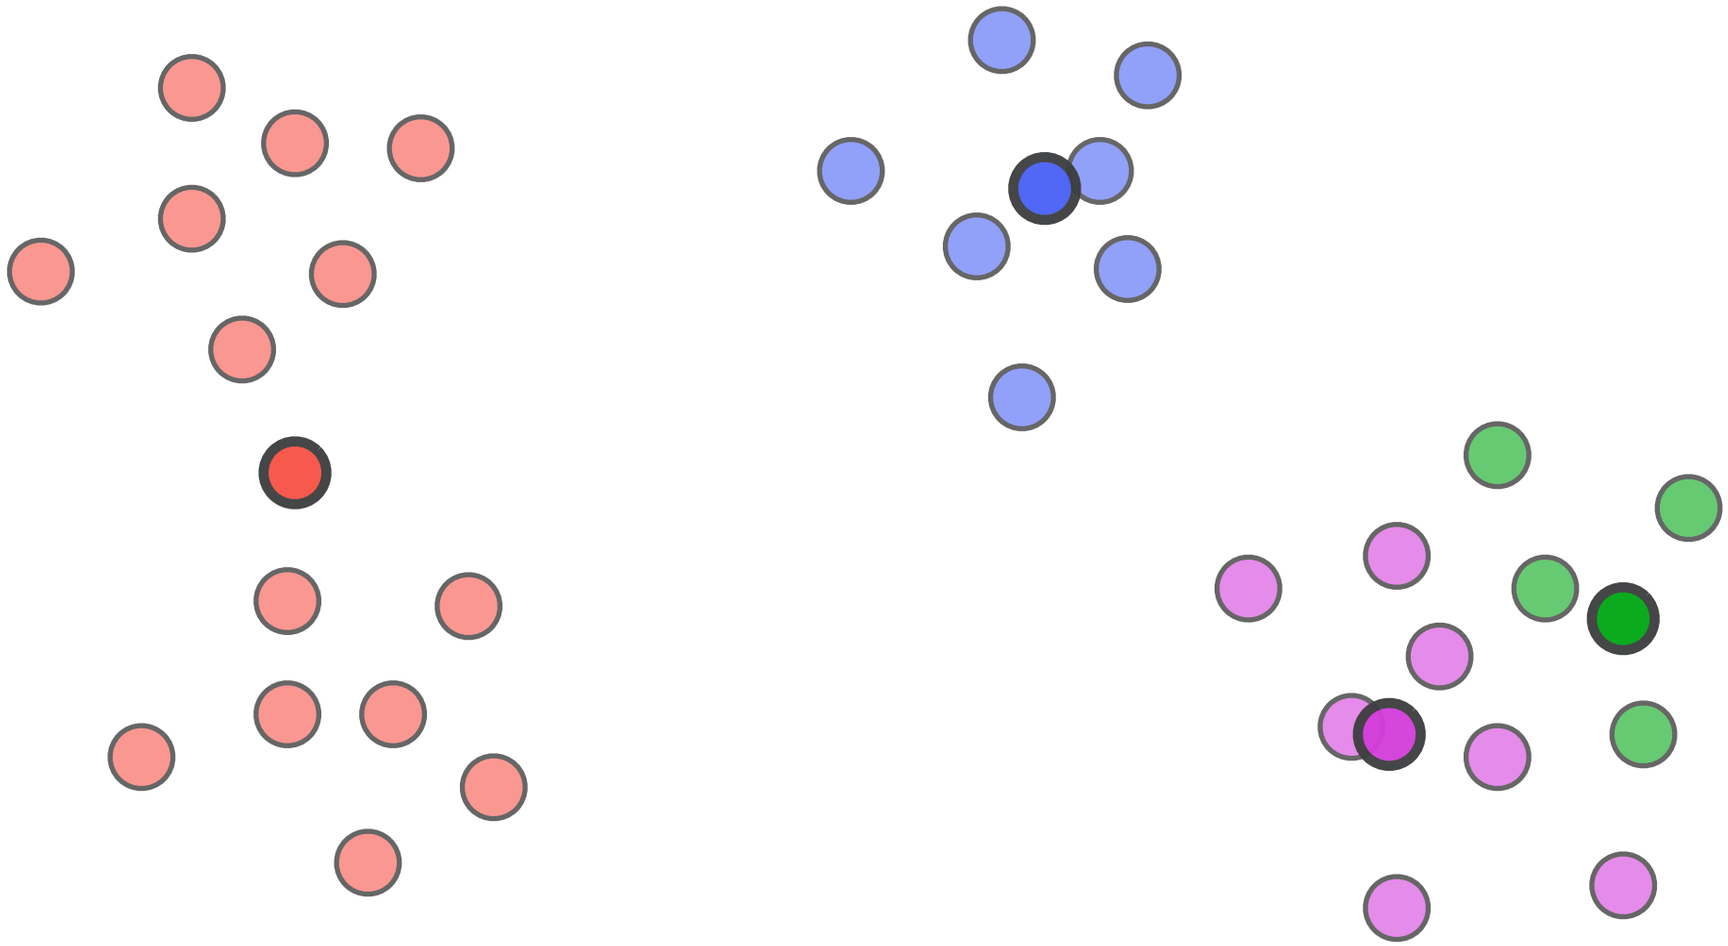

k-means algorithm: initialization with further-first traversal

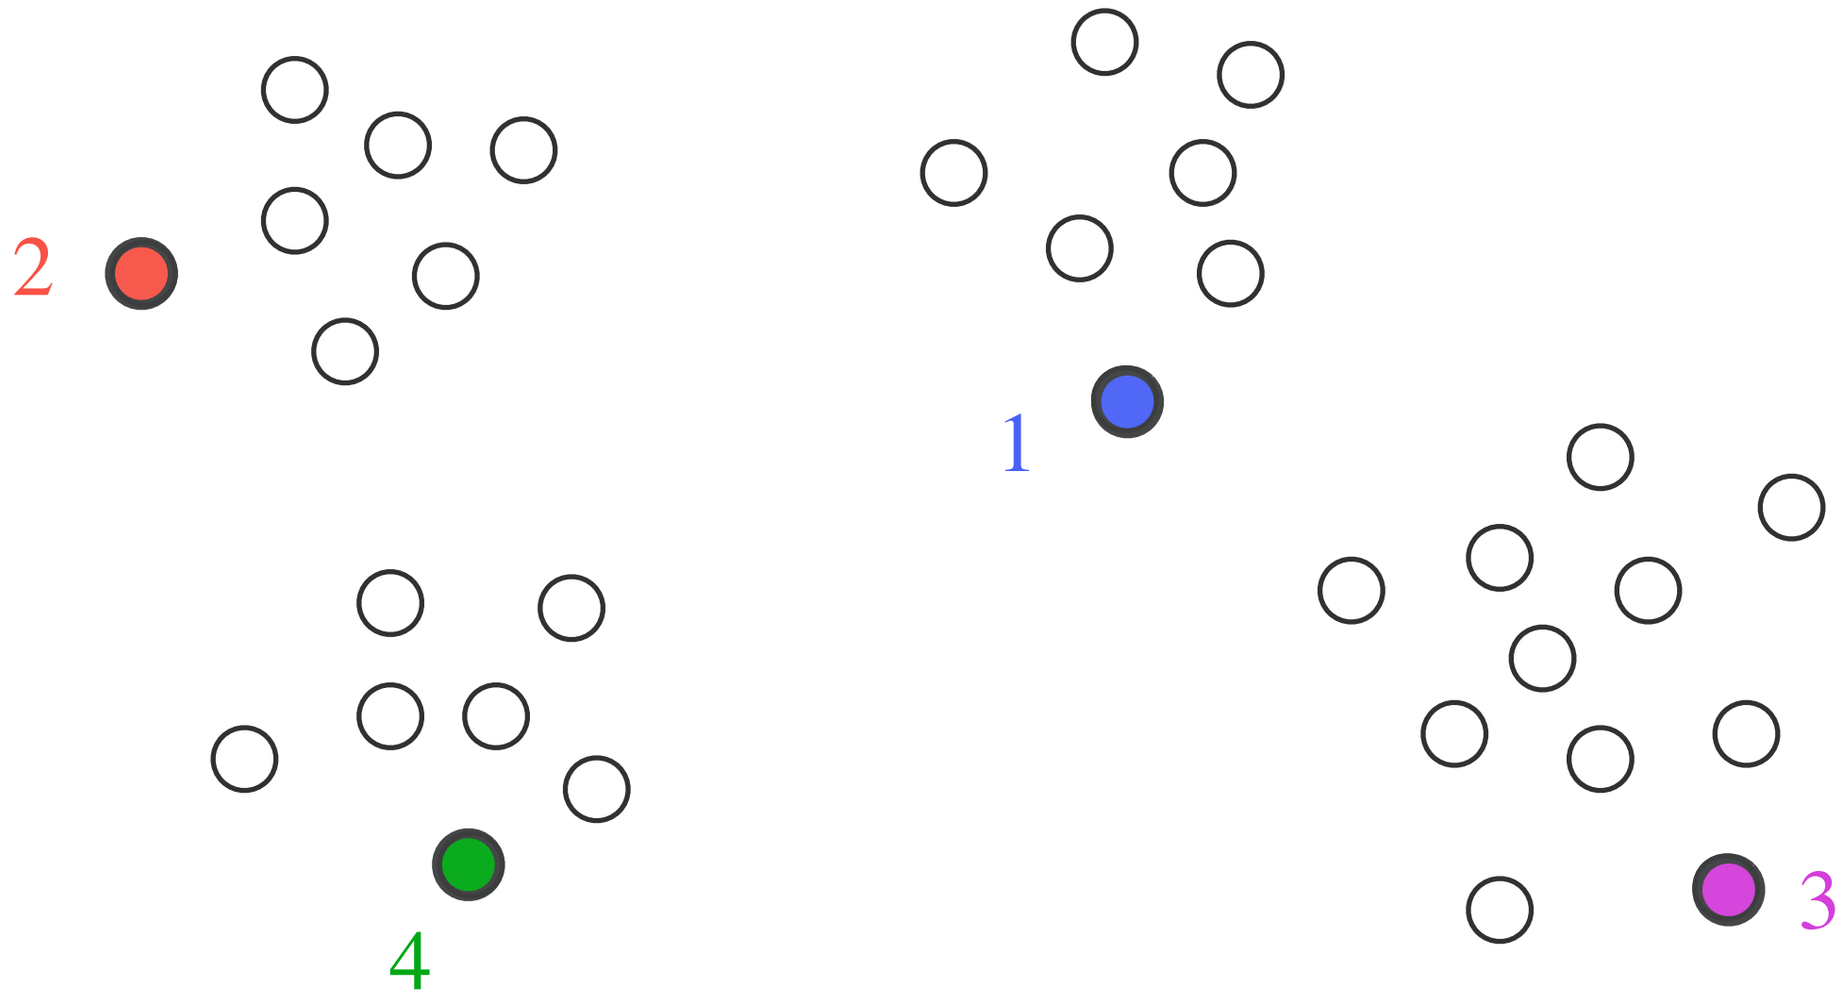

k-means algorithm: initialization with further-first traversal

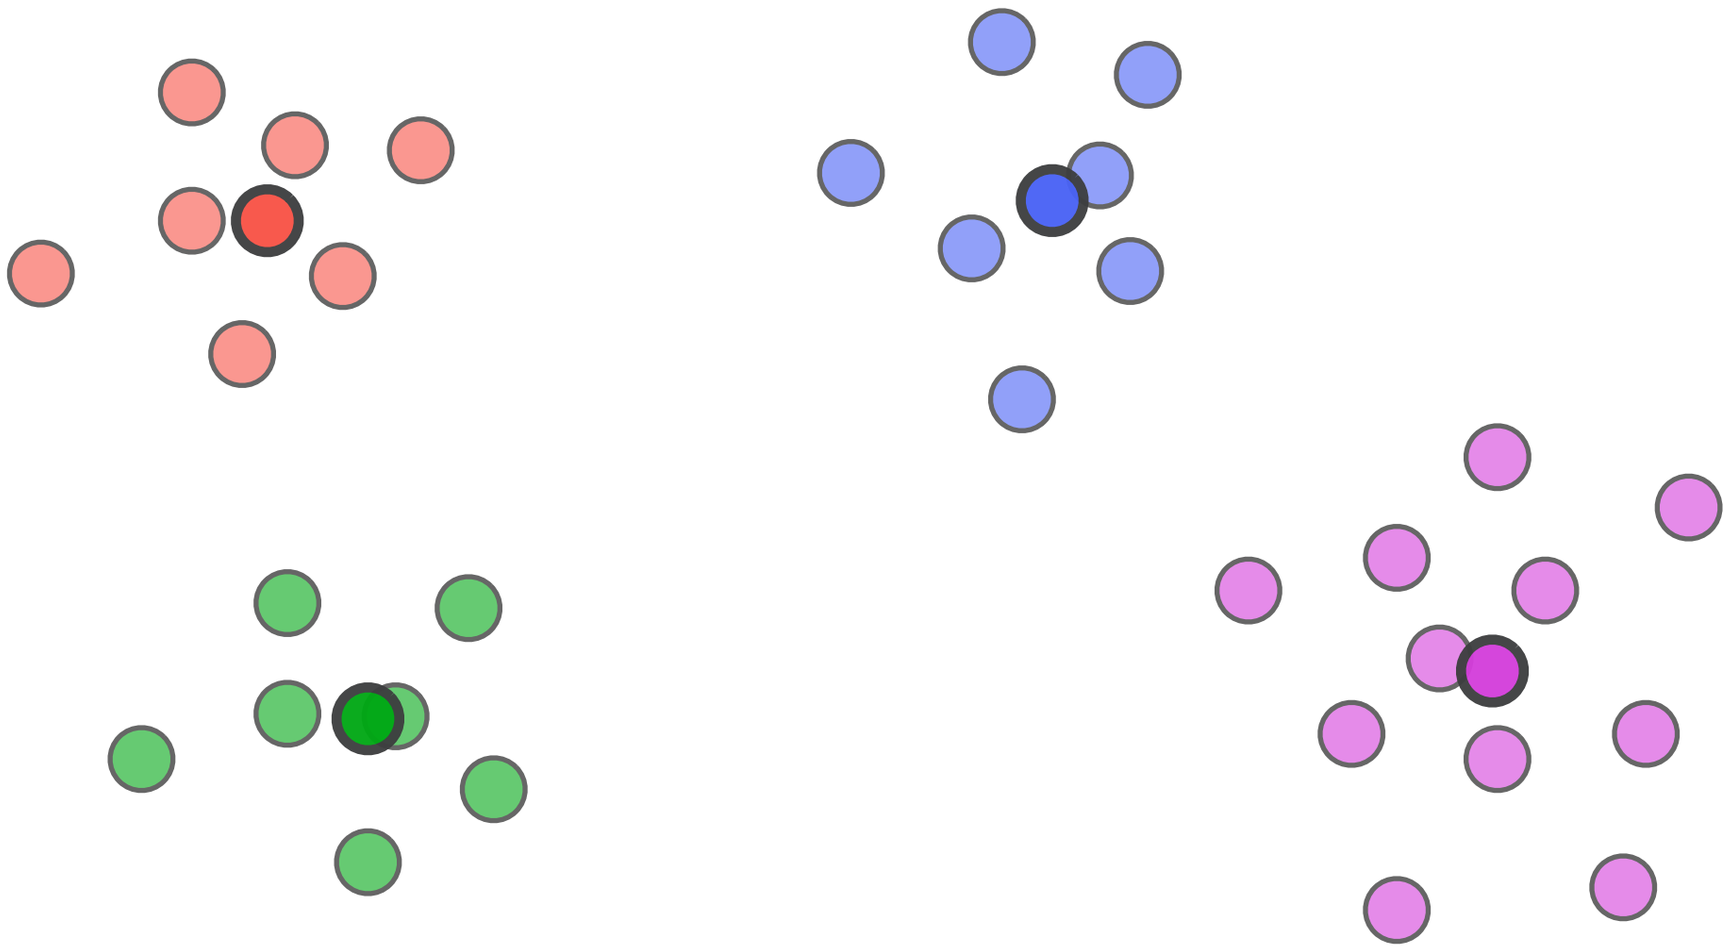

but... sensitive to outliers

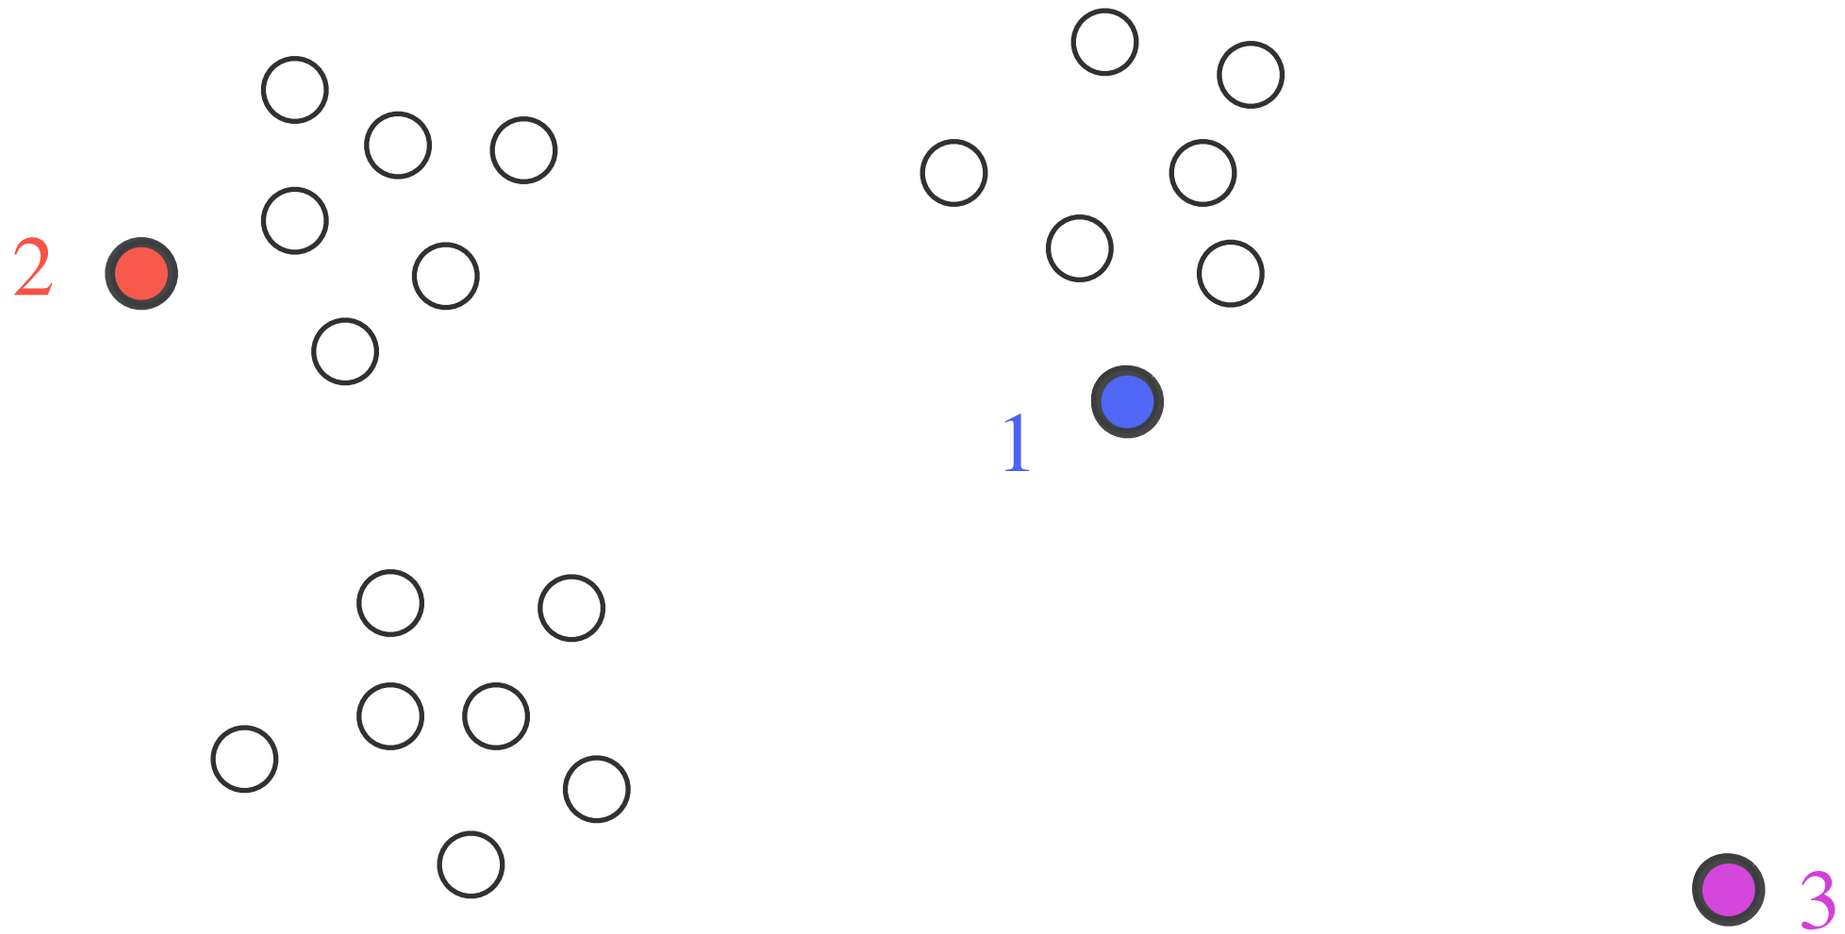

but... sensitive to outliers

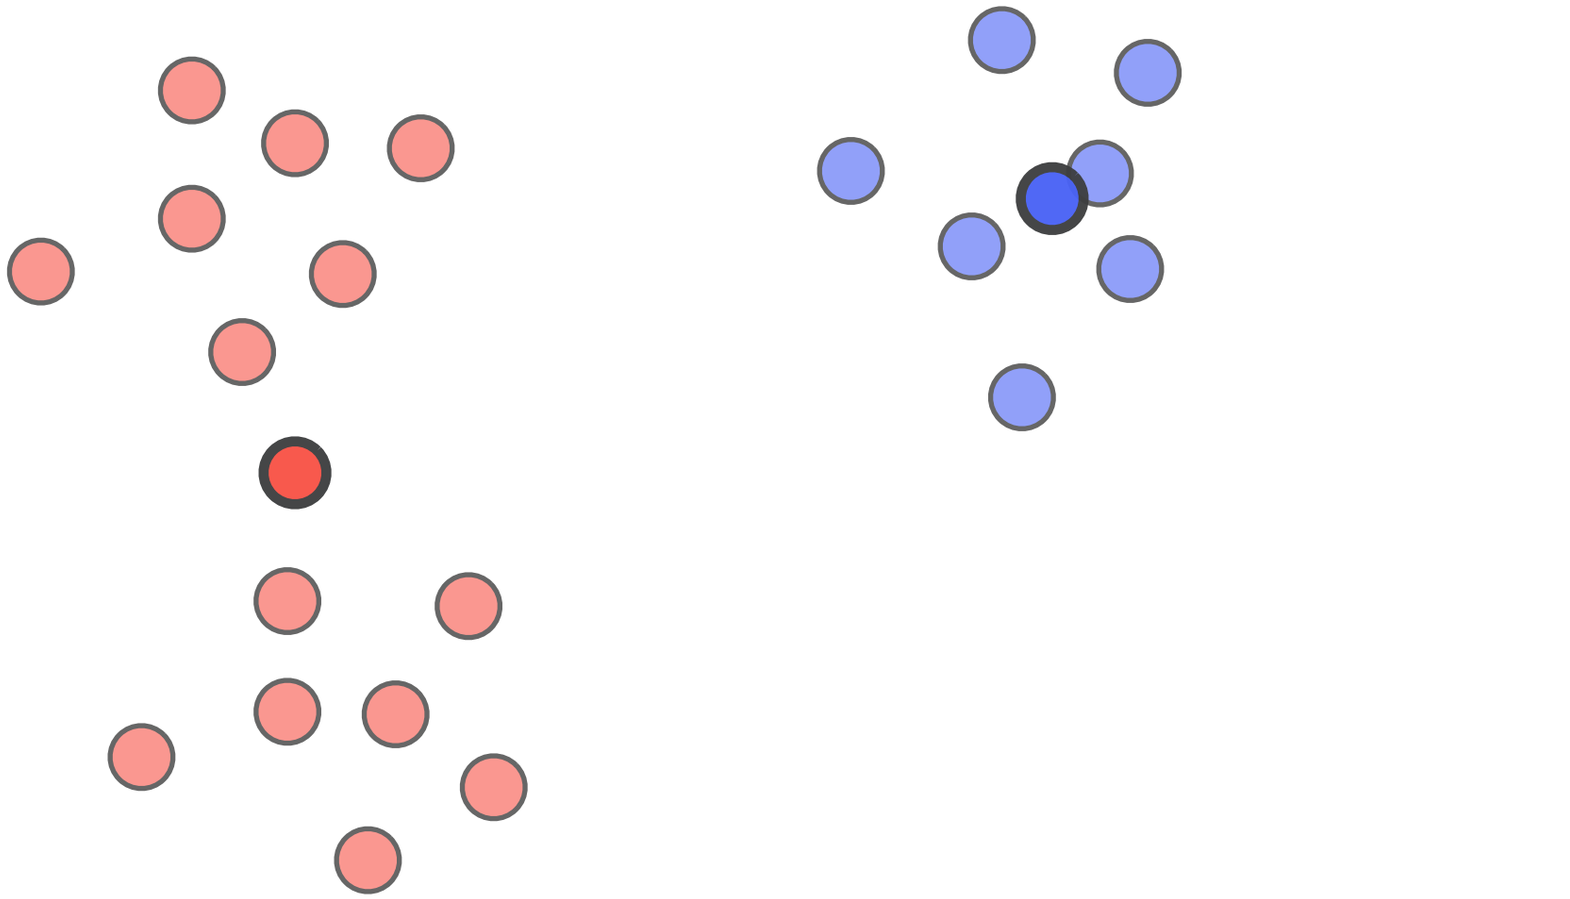

Here random may work well

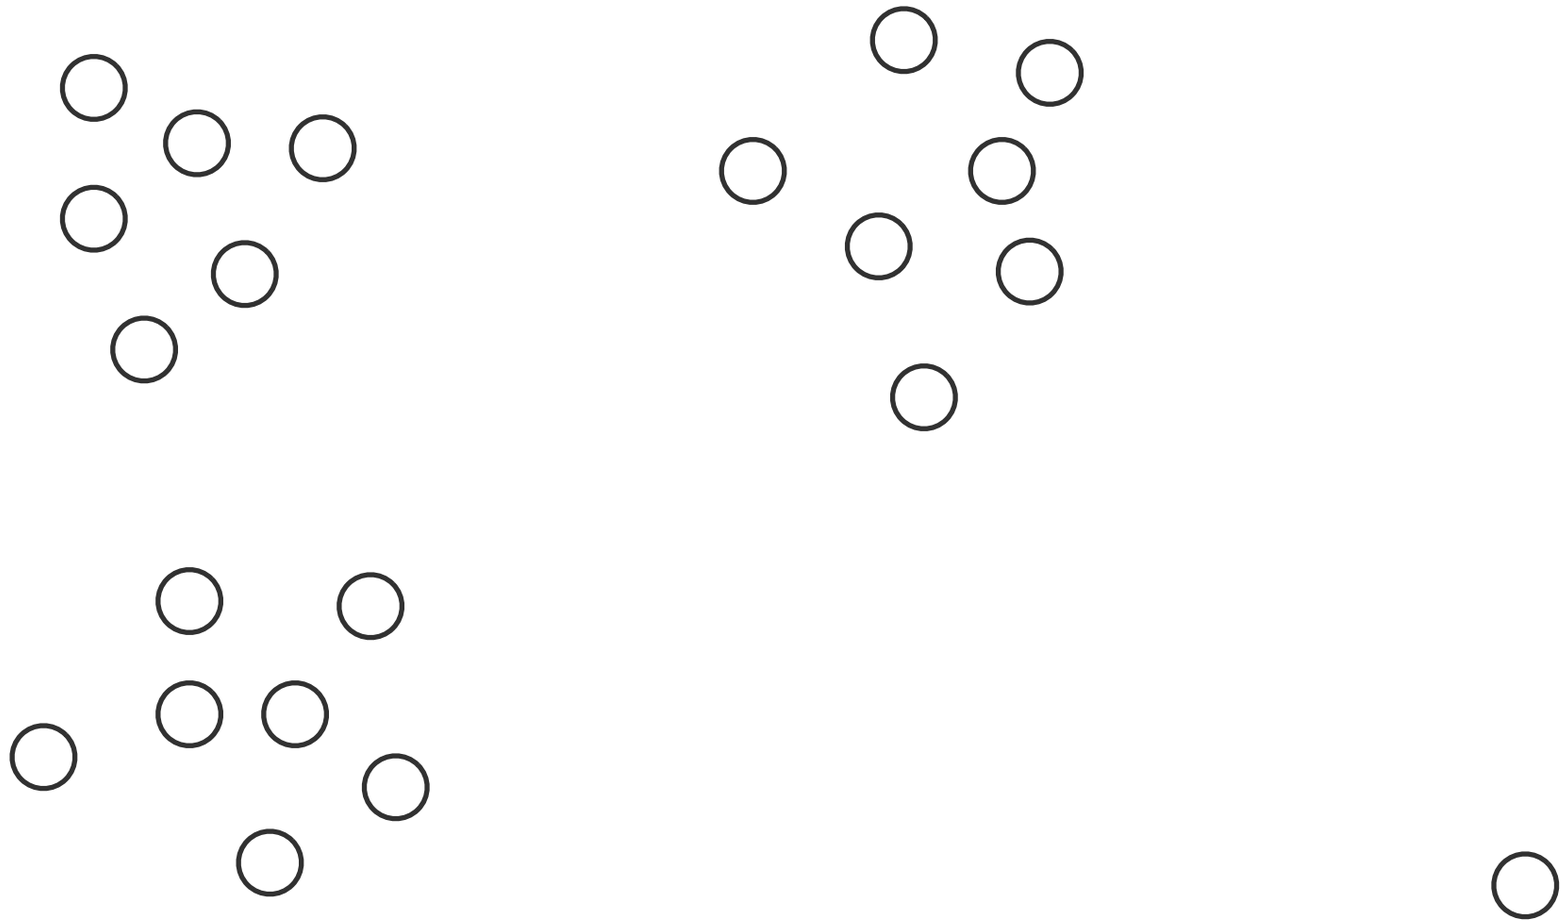

k-means++ algorithm

- interpolate between the two methods
- let D(x) be the distance between x and the nearest centers selected so far
- choose next center with probability proportional to

$$(D(x))^a = D^a(x)$$

- $a = 0$ random initialization
- $a = \infty$ furthest-first traversal
- $a = 2$ k-means++

k-means++ algorithm

- initialization phase:
  - choose the first center uniformly at random
  - choose next center with probability proportional to $D^2(x)$
- iteration phase:
  - iterate as in the k-means algorithm until convergence

k-means++ initialization

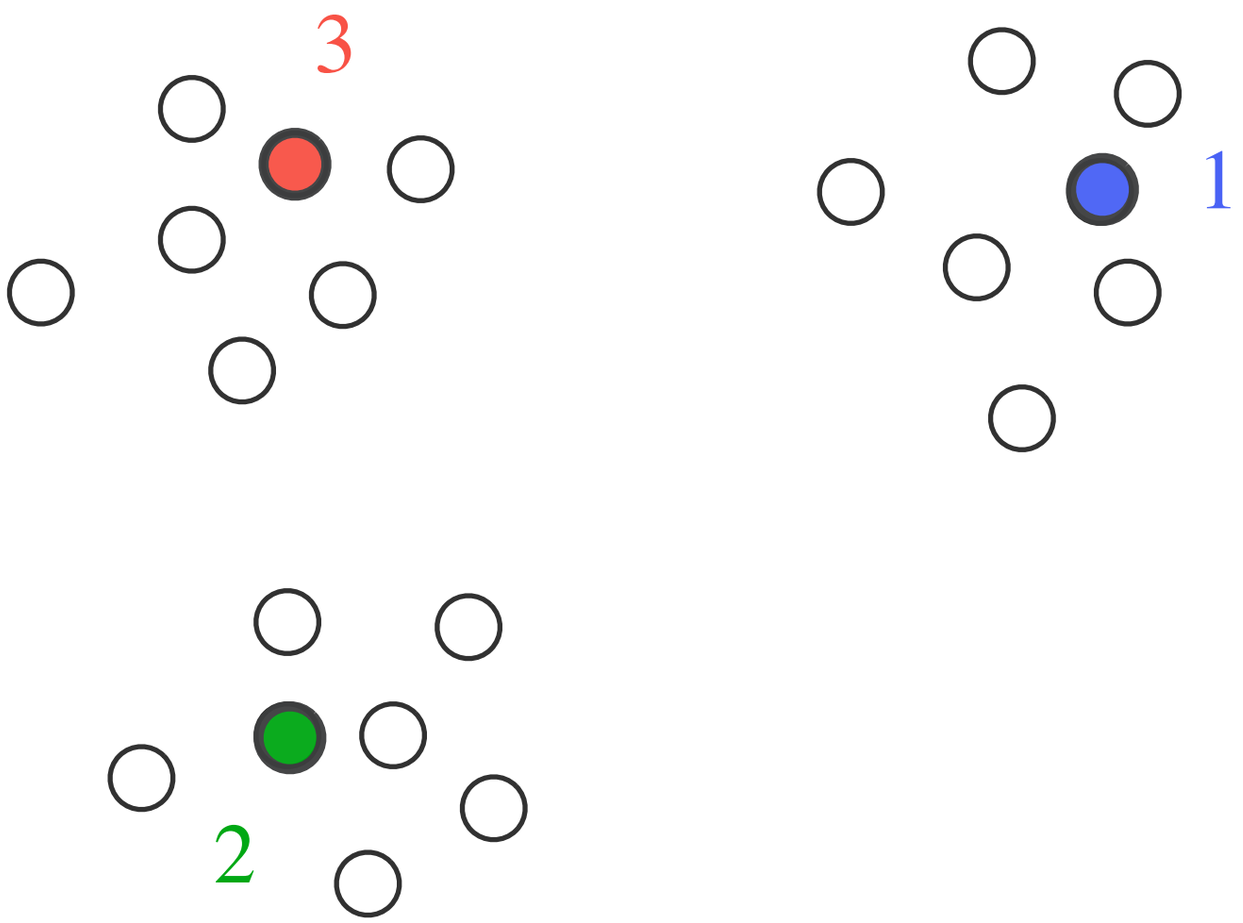

k-means++ result

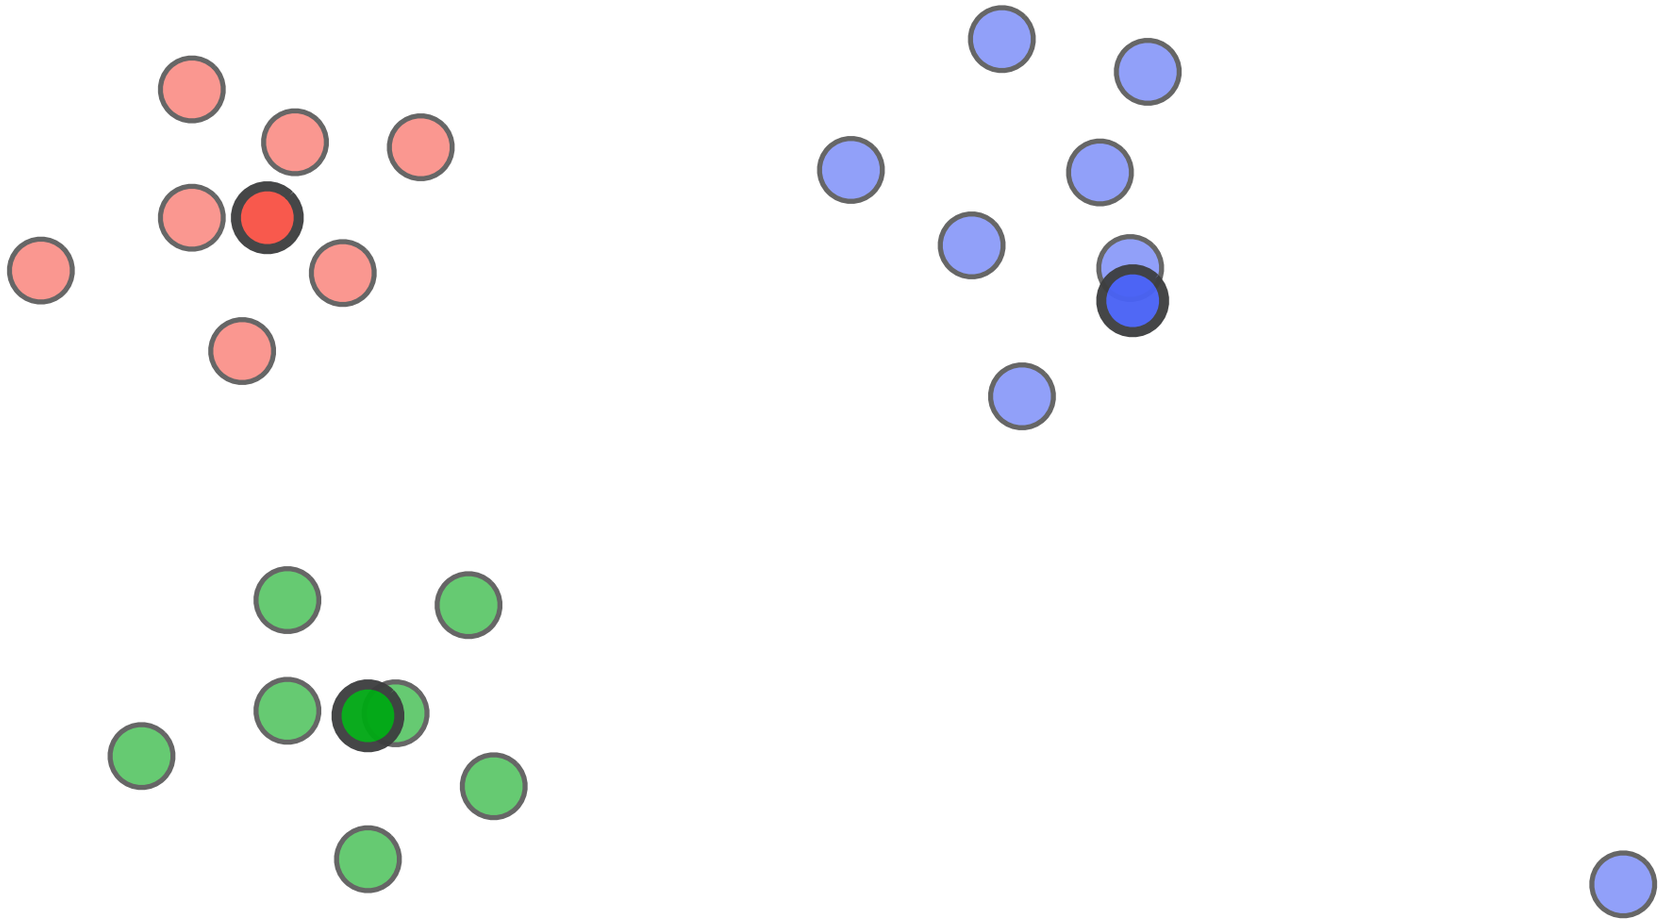

k-means++ provable guarantee

- approximation guarantee comes just from the first iteration (initialization)
- subsequent iterations can only improve cost

Lesson learned

- no reason to use k-means and not k-means++
- k-means++ :
  - easy to implement
  - provable guarantee
  - works well in practice

- Wiki: https://en.wikipedia.org/wiki/K-means%2B%2B
- Implemented in scikit-learn
- Even a parallel version!! k-means||

But.. how to choose the right k?

- Results depend on the choice of k
- How to choose the right k before even seen the data?
- Idea:
  - Iterate through different k and use a statistical test to find the most appropriate k!!
  - given that k-means tries to find spherical clusters… assume that clusters are spherical Gaussians!
  - use a statistical test to find if the models used fit the data better with the current k.